



Vadim Kubrak

Data Engineer & Data Analyst

Moscow, Russia — remote now; open to employer-supported relocation · kubrak15112006@gmail.com · t.me/Castiel68
archangel-gabri.github.io · github.com/Archangel-gabri · linkedin.com/in/vadimkubrak · ORCID 0009-0003-1553-2982

I work at **both ends of a data problem**: I build the pipelines that make the numbers available, and I do the analysis that makes them mean something. Python and SQL on PostgreSQL, ETL/ELT, relational modeling, scheduled jobs, FastAPI where an interface is needed — and I stay responsible for what I ship: Docker, Linux, CI/CD, TLS, monitoring, backups, incident response.

EXPERIENCE

Data Engineer — Nizhny Novgorod City Administration

Jun 2026 — present

- Data-engineering team building **ProcureCheck**: automated compliance checking of public procurement document packages under Russian federal laws 44-FZ and 223-FZ. Four documents go in; a short summary and a full report across twenty checks come out.
- Python, FastAPI, PDF and DOCX extraction (PyMuPDF, pdfplumber), Docker, GitLab CI. **More commits in the project than anyone else on the team.**
- Now in live use by a city department. Traced three false findings reported by its staff to a single root cause — a converter emitting text without word boundaries — and fixed the class rather than the three symptoms; separately found and closed a silent loss of 52 specification rows on a real document.

Data Analytics Intern — Gazprombank (JSC), Moscow

Mar — Aug 2026

- Pulled and joined data from several warehouse tables (transactions, customer profiles, operation history) in SQL and built an analytical view: customer segmentation, cohort retention, activity, average transaction value, churn.
- Presented findings as interactive Python dashboards (pandas, NumPy, Matplotlib/Plotly), each with a short business-readable takeaway rather than a wall of numbers.
- Built the internship's main deliverable: an internal FastAPI pipeline that collects data from several sources, computes the metrics and lets a language model draft the narrative commentary analysts previously wrote by hand. **Report preparation went from most of a working day to about fifteen minutes.**

Founder, Backend & Infrastructure Engineer — HubVPN

2024 — Aug 2026

- Built and operated a multi-node private network platform: a Remnawave control plane over an Xray/VLESS fleet with a German master node, a Moscow cascade and international exit nodes.
- Node provisioning through the Remnawave API, client-specific routing generation, Telegram-first onboarding on FastAPI, TLS, nginx, systemd, monitoring, backups and incident response — end to end. The platform was shut down cleanly in August 2026.

SELECTED PROJECTS

Ten repositories are public. Links go to code, not to claims.

Golf Club CRM — booking, membership and billing for a Moscow golf club, running in production for its client. **53 REST endpoints** on FastAPI (SQLAlchemy, Pydantic, JWT over httpOnly cookies), 25+ tables; React 18 / TypeScript / Vite / TanStack Query front end behind nginx with TLS; Docker Compose, Alembic, PostgreSQL. → [golf-crm](#)

LLM University Math Benchmark — open dataset and harness: 12 language models against 75 university mathematics problems, graded on the **solution process** rather than the final answer through an LLM-as-a-judge rubric; 887 of 900 responses scored, Friedman $\chi^2 = 351$, $p < 0.001$. Problems with reference solutions, raw model outputs, prompts and scoring code are public in Russian and English, CC BY 4.0 / MIT. → [dataset](#) · [harness](#)

Argus — local-first desktop command centre in Electron: own servers, subscriptions, finances and AI-provider quotas in one window. The store is SQLCipher with an Argon2id key and secrets never leave the privileged process. The interesting constraint was honesty of state — “unreachable” and “off” are different things, so one failed poll never marks a live machine as down. **994 tests across 87 files.** TypeScript, React, Tailwind, better-sqlite3. → [argus](#)

FX Analytics — currency-risk analysis for an exporting company (*project closed Aug 2026*). Asynchronous ingestion from Central Bank XML, MOEX ISS and yfinance, normalization, historical backfill, bulk upserts and scheduled reports; a four-agent pipeline on the Claude API with tool use (orchestrator → collector → quant + macro → reporter), plus an MCP server. About \$2.30 of model cost per full analysis.

MedSpravki REU — medical-certificate registry for the Physical Education department at Plekhanov REU. ASP.NET Core + PostgreSQL, with a local Ollama vision model reading scanned certificates into structured rows. Runs on real hardware. → [medspravki-reu](#)

Clean Listings — open-source browser extension, MIT. Filters resellers out of a classifieds search page by its own criterion rather than the site's. Profiled rather than guessed: **cut a 0.8-second render freeze to 0.19 seconds**; 28 regression tests, bilingual README. → [clean-listings](#)

Price Elsewhere — MV3 extension, no backend: open a product on one marketplace, see what it costs on two others. The marketplaces share no product identifier, so matching runs on brand, model and numbers — and it deliberately reports the *usual* price (a trimmed median) rather than the cheapest hit, because counterfeits and wrong bundles live in the cheap tail. When the match is uncertain it says so instead of promising a saving it cannot back. → [price-elsewhere](#)

Ghost — voice-and-screen desktop assistant in PySide6/Qt: hears a spoken question, reads what is on screen and streams the answer token by token into an overlay. A QThread owns the whole audio → speech-to-text → model chain and talks to the UI only through signals, so a two-second transcription never freezes the window it is typing into. Speech-to-text on Groq or fully offline through faster-whisper. → [ghost-assistant](#)

GORKYCODE 2025 (hackathon) — AI service for public-procurement automation: initial-price calculation and technical-specification generation under 44-FZ, multi-step requirement gathering over a webhook, GigaChat for analysing answers, pricing module on zakupki.gov.ru data. Specification drafting went from eight hours to minutes.

SKILLS

Data engineering: Python · SQL · ETL/ELT · async data ingestion · relational data modeling · data quality · scheduled jobs · historical backfill · bulk upsert · PostgreSQL · SQLite · SQLAlchemy · Alembic · openpyxl

Analytics: cohort retention · customer segmentation · churn and activity metrics · average transaction value · pandas · NumPy · Matplotlib · Plotly · statistics · Jupyter

Backend: FastAPI · Pydantic · REST · JWT · ASP.NET Core · Telegram Bot API

AI / LLM: Claude API · OpenAI API · OpenRouter · Ollama · MCP · RAG · multi-agent orchestration · prompt engineering · LLM-as-a-judge evaluation · embeddings and vector search · scikit-learn · PyTorch

Infrastructure & DevOps: Linux · Docker & Docker Compose · nginx · systemd · TLS · GitLab CI · GitHub Actions · CI/CD · Grafana · Prometheus · Playwright · Git · monitoring, backups and incident response on live nodes

Frontend: TypeScript · React 18 · Vite · TanStack Query · Tailwind CSS · Electron

EDUCATION & LANGUAGES

Plekhanov Russian University of Economics

BSc, Statistics / Data Science — in progress (2024 — 2028)

School 21 (Sber)

Data Science track — peer-to-peer C and algorithms (2025 — present)

Languages

Russian — native · English — C1 Advanced, [EF SET 89/100](#)

Research identity

ORCID [0009-0003-1553-2982](#)